

Midterm

Sining Leng

2025-03-24

Load data

```
load("dat1.RData")
load("dat2.RData")

dat1 =
  dat1 |>
  select(-id)

dat2 =
  dat2 |>
  select(-id)
```

Explorey Data Analysis and Visualization

Summary statistics

dat1

```
summarize_cat = function(data, var_name) {

  var_label = deparse(substitute(var_name))

  out_df =
    data |>
    count({{ var_name }}) |>
    mutate(
      Percent = round(100 * n / sum(n), 1),
      Variable = var_label,
      Level = as.character({{ var_name }})
    ) |>
    rename(N = n) |>
    select(Variable, Level, N, Percent)

  return(out_df)
}

gender <- summarize_cat(dat1, gender)
race <- summarize_cat(dat1, race)
smoking <- summarize_cat(dat1, smoking)
```

```
diabetes <- summarize_cat(dat1, diabetes)
hypertension <- summarize_cat(dat1, hypertension)
bind_rows(gender, race, smoking, diabetes, hypertension) %>%
  knitr::kable()
```

Variable	Level	N	Percent
gender	0	2573	51.5
gender	1	2427	48.5
race	1	3221	64.4
race	2	278	5.6
race	3	1036	20.7
race	4	465	9.3
smoking	0	3010	60.2
smoking	1	1504	30.1
smoking	2	486	9.7
diabetes	0	4228	84.6
diabetes	1	772	15.4
hypertension	0	2702	54.0
hypertension	1	2298	46.0

```
summarize_cont = function(data, var_name) {

  var_label = deparse(substitute(var_name))

  out_df =
    data |>
    summarize(
      Variable = var_label,
      Median = round(median({{ var_name }}), na.rm = TRUE), 1),
      Q1 = round(quantile({{ var_name }}), 0.25, na.rm = TRUE), 1),
      Q3 = round(quantile({{ var_name }}), 0.75, na.rm = TRUE), 1)
    ) |>
    mutate(
      IQR = paste0("[", Q1, ", ", Q3, "]")
    ) |>
    select(Variable, Median, IQR)

  return(out_df)
}

age = summarize_cont(dat1, age)
bmi = summarize_cont(dat1, bmi)
height = summarize_cont(dat1, height)
weight = summarize_cont(dat1, weight)
SBP = summarize_cont(dat1, SBP)
LDL = summarize_cont(dat1, LDL)
time = summarize_cont(dat1, time)
log_anti = summarize_cont(dat1, log_antibody)
bind_rows(age, bmi, height, weight, SBP, LDL, time, log_anti) %>%
  knitr::kable()
```

Variable	Median	IQR
age	60.0	[57, 63]
bmi	27.6	[25.8, 29.5]
height	170.1	[166.1, 174.2]
weight	80.1	[75.4, 84.9]
SBP	130.0	[124, 135]
LDL	110.0	[96, 124]
time	106.0	[76, 138]
log_antibody	10.1	[9.7, 10.5]

dat2

```
gender <- summarize_cat(dat2, gender)
race <- summarize_cat(dat2, race)
smoking <- summarize_cat(dat2, smoking)
diabetes <- summarize_cat(dat2, diabetes)
hypertension <- summarize_cat(dat2, hypertension)
bind_rows(gender, race, smoking, diabetes, hypertension) %>%
  knitr::kable()
```

Variable	Level	N	Percent
gender	0	509	50.9
gender	1	491	49.1
race	1	663	66.3
race	2	55	5.5
race	3	199	19.9
race	4	83	8.3
smoking	0	601	60.1
smoking	1	296	29.6
smoking	2	103	10.3
diabetes	0	843	84.3
diabetes	1	157	15.7
hypertension	0	544	54.4
hypertension	1	456	45.6


```
age = summarize_cont(dat2, age)
bmi = summarize_cont(dat2, bmi)
height = summarize_cont(dat2, height)
weight = summarize_cont(dat2, weight)
SBP = summarize_cont(dat2, SBP)
LDL = summarize_cont(dat2, LDL)
time = summarize_cont(dat2, time)
log_anti = summarize_cont(dat2, log_antibody)
bind_rows(age, bmi, height, weight, SBP, LDL, time, log_anti) %>%
  knitr::kable()
```

Variable	Median	IQR
age	60.0	[57, 63]
bmi	27.6	[25.8, 29.6]

Variable	Median	IQR
height	170.2	[166.1, 174.2]
weight	80.2	[75.3, 84.4]
SBP	130.0	[124, 135]
LDL	112.0	[96, 124]
time	171.0	[140, 205]
log_antibody	9.9	[9.5, 10.3]

Distribution of antibody levels

```
ggplot(dat1, aes(x = log_antibody)) +
  geom_histogram(binwidth = 0.2, fill = 'skyblue', color = 'black') +
  labs(x = "Log-transformed antibody level",
       y = "Frequency",
       title = "Distribution of Log-Transformed Antibody Levels",
       subtitle = "dat1") +
  theme_classic()
```




```
ggplot(dat2, aes(x = log_antibody)) +
  geom_histogram(binwidth = 0.2, fill = 'skyblue', color = 'black') +
  labs(x = "Log-transformed antibody level",
       y = "Frequency",
```

```

title = "Distribution of Log-Transformed Antibody Levels",
subtitle = "dat2") +
theme_classic()

```



Scatterplots of Continuous Predictors vs. Log-Antibody


```

dat1 %>%
  select(age, height, weight, bmi, SBP, LDL, time, log_antibody) %>%
  pivot_longer(
    cols = c(age, height, weight, bmi, SBP, LDL, time),
    names_to = "predictor",
    values_to = "value"
  ) %>%
  ggplot(aes(x = value, y = log_antibody)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", se = FALSE, color = "blue") +
  facet_wrap(~ predictor, scales = "free_x") +
  labs(
    x = "Predictor Value",
    y = "Log-Antibody",
    title = "Scatterplots of Continuous Predictors vs. Log-Antibody"
  ) + theme_test()

```

```
## 'geom_smooth()' using formula = 'y ~ x'
```


Scatterplots of Continuous Predictors vs. Log-Antibody



Correlation Matrix of Continuous Variables

```
dat1 %>%
  select(age, height, weight, bmi, SBP, LDL, time, log_antibody) %>%
  cor(use = "complete.obs") %>%
  corrrplot(type = "full",
            title = "Correlation Matrix of Continuous Variables",
            addCoef.col = "black",
            mar = c(0,0,2,0))
```

Correlation Matrix of Continuous Variables



Boxplots of log_antibody by categorical variables

```

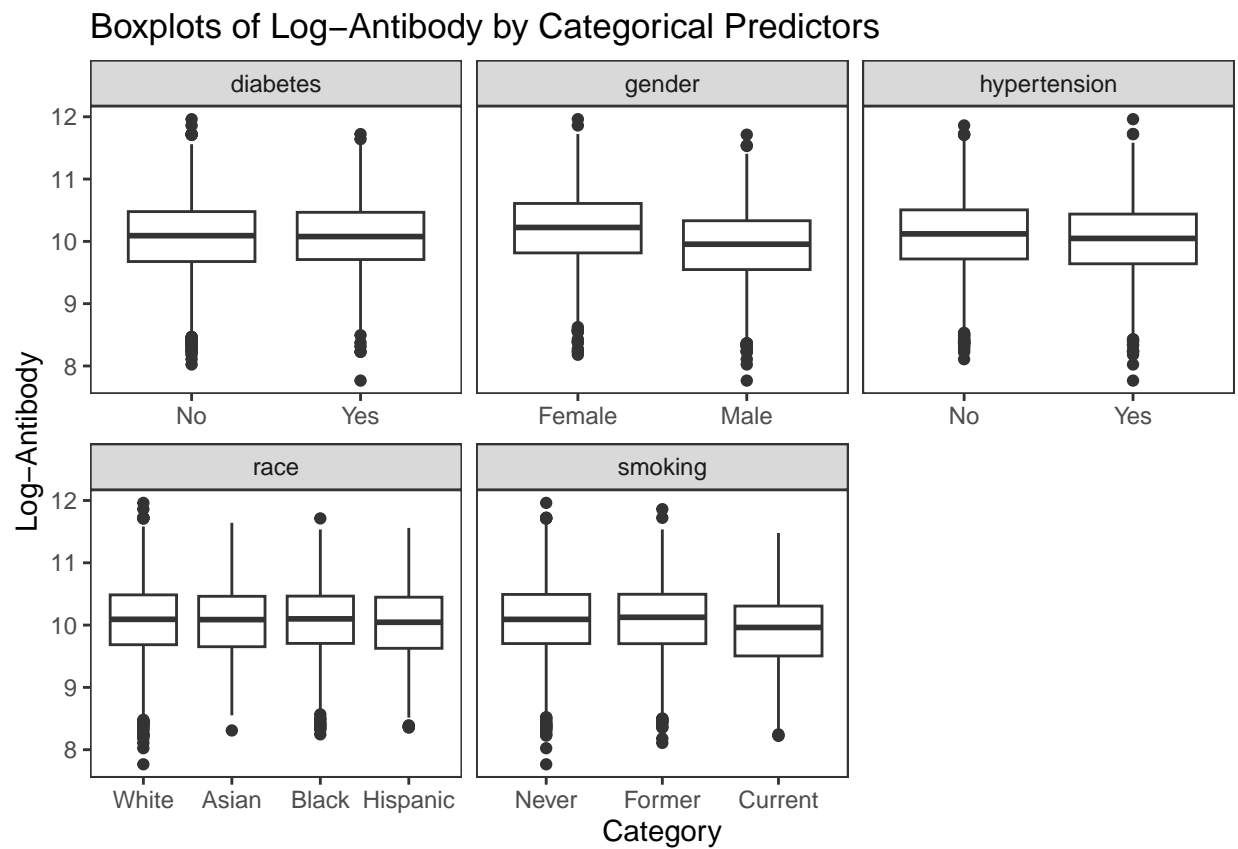
dat1 %>%
  select(log_antibody, gender, race, smoking, diabetes, hypertension) %>%
  mutate(
    gender = factor(gender,
      levels = c(0, 1),
      labels = c("Female", "Male")),
    race = factor(race,
      levels = c(1, 2, 3, 4),
      labels = c("White", "Asian", "Black", "Hispanic")),
    smoking = factor(smoking,
      levels = c(0, 1, 2),
      labels = c("Never", "Former", "Current")),
    diabetes = factor(diabetes,
      levels = c(0, 1),
      labels = c("No", "Yes")),
    hypertension = factor(hypertension,
      levels = c(0, 1),
      labels = c("No", "Yes"))
  ) %>%
  pivot_longer(
    cols = c(gender, race, smoking, diabetes, hypertension),
    names_to = "predictor",

```

```

  values_to = "category"
) %>%
ggplot(aes(x = category, y = log_antibody)) +
geom_boxplot() +
facet_wrap(~ predictor, scales = "free_x") +
labs(
  x = "Category",
  y = "Log-Antibody",
  title = "Boxplots of Log-Antibody by Categorical Predictors"
) + theme_test()

```



Antibody level and time

```

p1 <- ggplot(dat1, aes(x = time, y = log_antibody)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", color = "blue") +
  labs(title = "Antibody Level vs Time Since Vaccination (dat1)") +
  theme_classic()

p2 <- ggplot(dat2, aes(x = time, y = log_antibody)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess", color = "blue") +



```



```
labs(title = "Antibody Level vs Time Since Vaccination (dat2)") +  
theme_classic()
```

```
p1+p2
```

```
## 'geom_smooth()' using formula = 'y ~ x'  
## 'geom_smooth()' using formula = 'y ~ x'
```



Models Building

cross-validation

```
set.seed(123)  
ctrl = trainControl(method = "cv", number = 10)
```

LASSO

```
set.seed(123)  
model_lasso <- train(  
  log_antibody ~ .,
```

```

data = dat1,
method = "glmnet",
trControl = ctrl,
tuneGrid = expand.grid(alpha = 1,
                        lambda = exp(seq(2, -30, length = 100)))
))

```

```


## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo,
## : There were missing values in resampled performance measures.

```

```

plot(model_lasso, xTrans = log)

```



```

model_lasso$bestTune

```

```

##      alpha      lambda
## 60      1 1.792538e-05

```

GAM

```

set.seed(123)
model_gam <- train(

```

```

log_antibody ~ .,
data = dat1,
method = "gam",
trControl = ctrl
)

```

```
model_gam$finalModel
```

```

##
## Family: gaussian
## Link function: identity
##
## Formula:
## .outcome ~ gender + race2 + race3 + race4 + smoking1 + smoking2 +
##      diabetes + hypertension + s(age) + s(SBP) + s(LDL) + s(bmi) +
##      s(time) + s(height) + s(weight)
##
## Estimated degrees of freedom:
## 0.991 0.000 0.000 4.179 7.892 1.234 0.000
## total = 23.3
##
## GCV score: 0.2786734

```

```
model_gam$bestTune
```

```





## select method
## 2 TRUE GCV.Cp

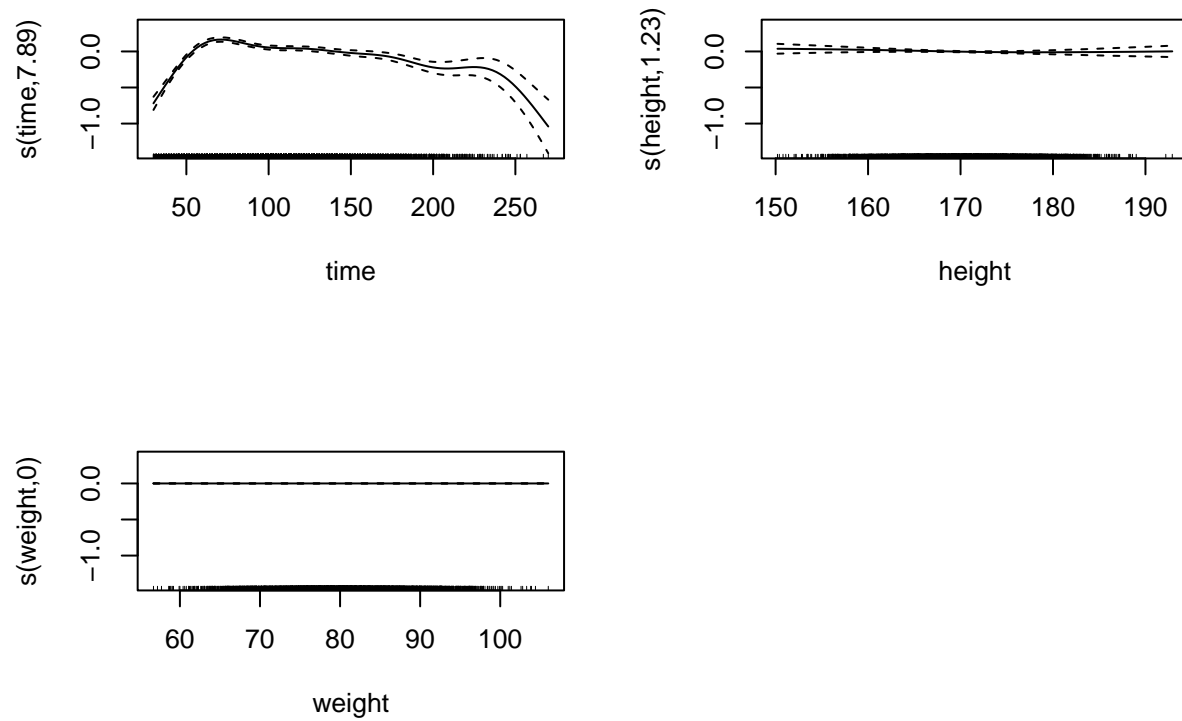
```

```

par(mfrow = c(2,2))
plot(model_gam$finalModel)

```






MARS

```
set.seed(123)
model_mars <- train(
  log_antibody ~ .,
  data = dat1,
  method = "earth",
  trControl = ctrl,
  tuneLength = 5
)


plot(model_mars)
```




```
coef(model_mars$finalModel)
```

```
## (Intercept) h(27.8-bmi) h(time-57) h(57-time) gender h(age-59)
## 10.847446930 -0.061997354 -0.002254182 -0.033529326 -0.296290451 -0.022957648
## h(59-age) smoking2 h(bmi-23.7)
## 0.016138468 -0.205126851 -0.084380175
```

```
vip(model_mars$finalModel, type = "nsubsets")
```



```
vip(model_mars$finalModel, type = "rss")
```



Predictions and Model Evaluation

```
set.seed(123)

pred_lasso <- predict(model_lasso, newdata = dat2)
pred_mars <- predict(model_mars, newdata = dat2)
pred_gam <- predict(model_gam, newdata = dat2)

resample = resamples(list(lasso = model_lasso, gam = model_gam, mars = model_mars))
summary(resample)
```

```
##
## Call:
## summary.resamples(object = resample)
##
## Models: lasso, gam, mars
## Number of resamples: 10
##
## MAE
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## lasso 0.4262302 0.4320751 0.4340211 0.4390288 0.4427211 0.4724617    0
## gam   0.4078470 0.4112544 0.4168810 0.4225034 0.4281610 0.4635896    0
## mars  0.4078837 0.4098168 0.4163554 0.4220725 0.4285910 0.4641090    0
```



```
##
## RMSE
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## lasso 0.5327885 0.5395192 0.5514583 0.5507990 0.5562623 0.5834205    0
## gam   0.5055708 0.5150506 0.5243183 0.5277724 0.5349189 0.5715483    0
## mars  0.5048757 0.5121685 0.5240607 0.5274715 0.5361255 0.5724126    0
##
## Rsquared
##      Min.    1st Qu.    Median      Mean    3rd Qu.      Max. NA's
## lasso 0.1001182 0.1324550 0.1530212 0.1478575 0.1617642 0.1910705    0
## gam   0.1586929 0.1935757 0.2259290 0.2177415 0.2449893 0.2595686    0
## mars  0.1566346 0.1954579 0.2274358 0.2187657 0.2465799 0.2627660    0
```

MARS model has the lowest RMSE (0.5276)